

Maximum Marks: 70

Time Allowed: 3 Hours

Note:- Attempt all questions. Internal choice is given.

SECTION – A

I. Short Answer Type Questions (Attempt Any Five) [5×4=20]

- i. Define Natural Language Processing (NLP) and list two real-world applications.
- ii. What is the role of Named Entity Recognition (NER) in NLP tasks?
- iii. Explain the concept of word embeddings with reference to Word2Vec and GloVe.
- iv. Differentiate between BERT and GPT in terms of language modeling approaches.
- v. What are the key steps in the NLP project pipeline?
- vi. Define Self-Attention and explain its significance in transformer models.
- vii. What are deepfakes? Mention any two safety measures to detect or prevent them.

SECTION – B

II. Long Answer Type Questions (Attempt Any Four) [4×10=40]

- i. Discuss the various **linguistic foundations** of NLP, such as syntax, semantics, ambiguity, and code-switching, with suitable examples.
- ii. Explain **data preprocessing techniques** in NLP. How do tokenization, stopword removal, and normalization improve model performance?
- iii. Describe the core components of the **Transformer architecture**. How do encoder-decoder mechanisms function?
- iv. Explain **sampling strategies** used in text generation (Greedy, Beam, Top-k, Top-p). Which strategies improve diversity and fluency?
- v. Analyze the challenges and risks associated with **Cross-Modal Generative AI**, including synthetic media and misuse. Mention policy-level safety measures.
- vi. Compare **Contextual Embeddings** (e.g., BERT, ELMo) with traditional word embeddings (e.g., Word2Vec). How does context influence meaning?